

K8s Nemesis

Requirements and analysis model

Product description

Product Description: A cloud-native service designed to enhance application scalability and performance in Kubernetes (K8s) environments, specifically targeting machine learning (ML) applications with unique scaling and resource requirements.

Team: Sergey Lokhmatikov, Andrey Tamplon, Roman Kuzmenko, Vasiliy Lopatkin

Repo: https://github.com/Lokhmat/k8s_nemesis

Report: https://github.com/Lokhmat/k8s_nemesis/blob/main/final_slides.pdf


MLOps engineer

Background: Focuses on deploying, monitoring, and maintaining ML models in production. He's skilled in cloud platforms, CI/CD pipelines, and container tools like Docker and Kubernetes.

Needs: Tools that help with continuous integration and deployment of ML models make it easier to quickly update and improve models without interrupting current services.

Challenges: Making sure models are reliable, keeping consistency across environments, and handling complex ML pipelines and dependencies.

Goal: Automate deployment and monitoring for optimal model performance.





Background: Engine

Needs: Elegant and customizable tool for managing k8_s cluster. Tool that can reduce cost of maintenance of a large scale architecture while ensuring high SLA and good performance in a complicated systems.

Challenges: Balancing infrastructure maintenance cost vs service uptime. Managing distributed cluster, and event driven scaling based on customizable set of rules and metrics. Ensuring security and stability of a system.

Goal: Deployment and maintenance of a reliable distributed system.
Increasing usage efficiency of a current organisation's cloud architecture.

Story Map

Storymap

Persona goal

Equal load
distribution
across cluster

Save money
and cluster
resources

Integrate
with k8s
cluster

Persona task

Ensure
minimal
latency

Guarantee
high level
availability of
cluster

Count of
pods
depends on
current load

Fair requests
proxying

Stop using
resources if
application is
not processing
requests

Allow to define
the most
suitable
settings for
k8s cluster

Activate
application if it
needs to handle
requests but not
active

Integration
with k8s
cluster

Story Map

Product feature

Automatically
scale cluster
up

Request
proxying with
different
policies

Automatically
scale cluster
down

Disable
application if
there are no
requests

Activate
application if
request comes
in and system
is disabled

Define scalling
rules based on
cluster load
metrics

Define
proxying rules
based on
cluster load
metrics

Fetch k8s
cluster
metrics

Monitor
current pods
liveliness
state

Stories

Increase
cluster pods
amount

Proxy
requests

Decrease
cluster pods
amount

Disable
cluster

Activate
cluster

Set
activation/de
activation
rule

Set load
balancing
policy

Get
CPU/RAM/DIS
K metrics
from k8s

Get current
pods
liveliness
state

Set load
balancing
policy

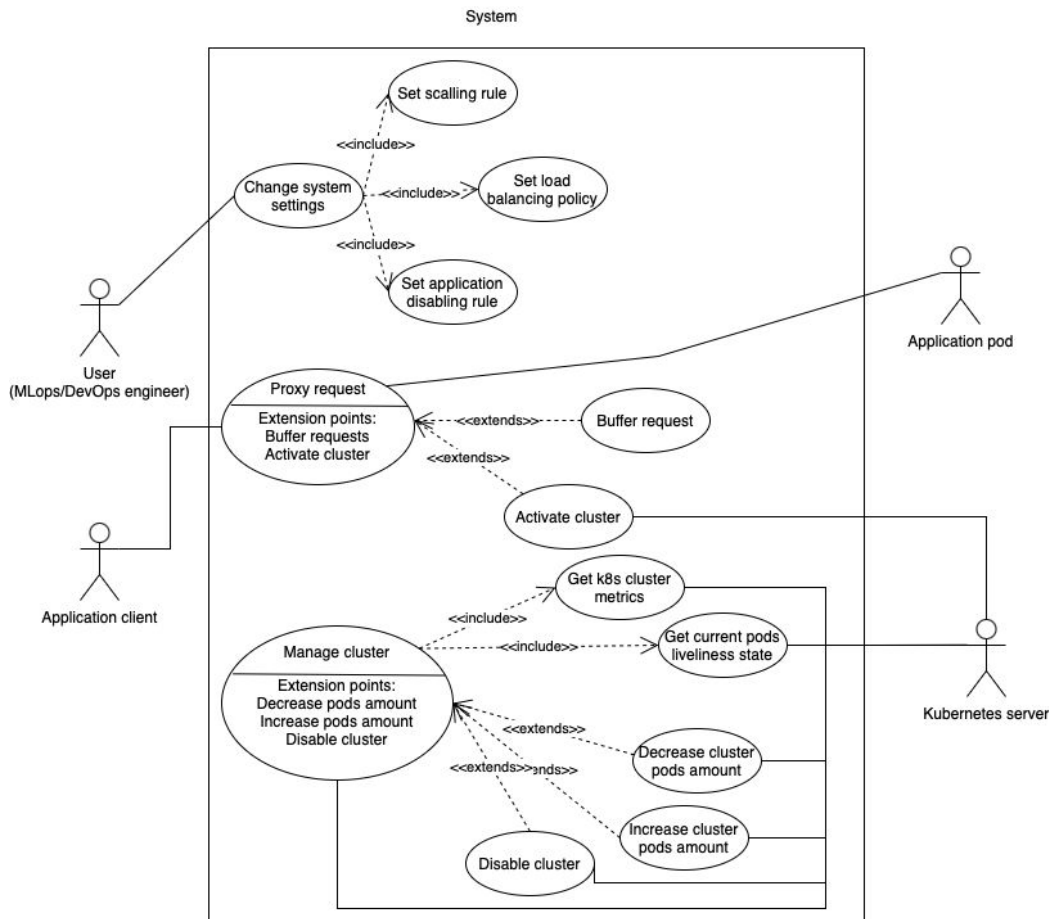
Buffer
request

Set service
scalling rule

Use case diagram

Textual use case scenarios:

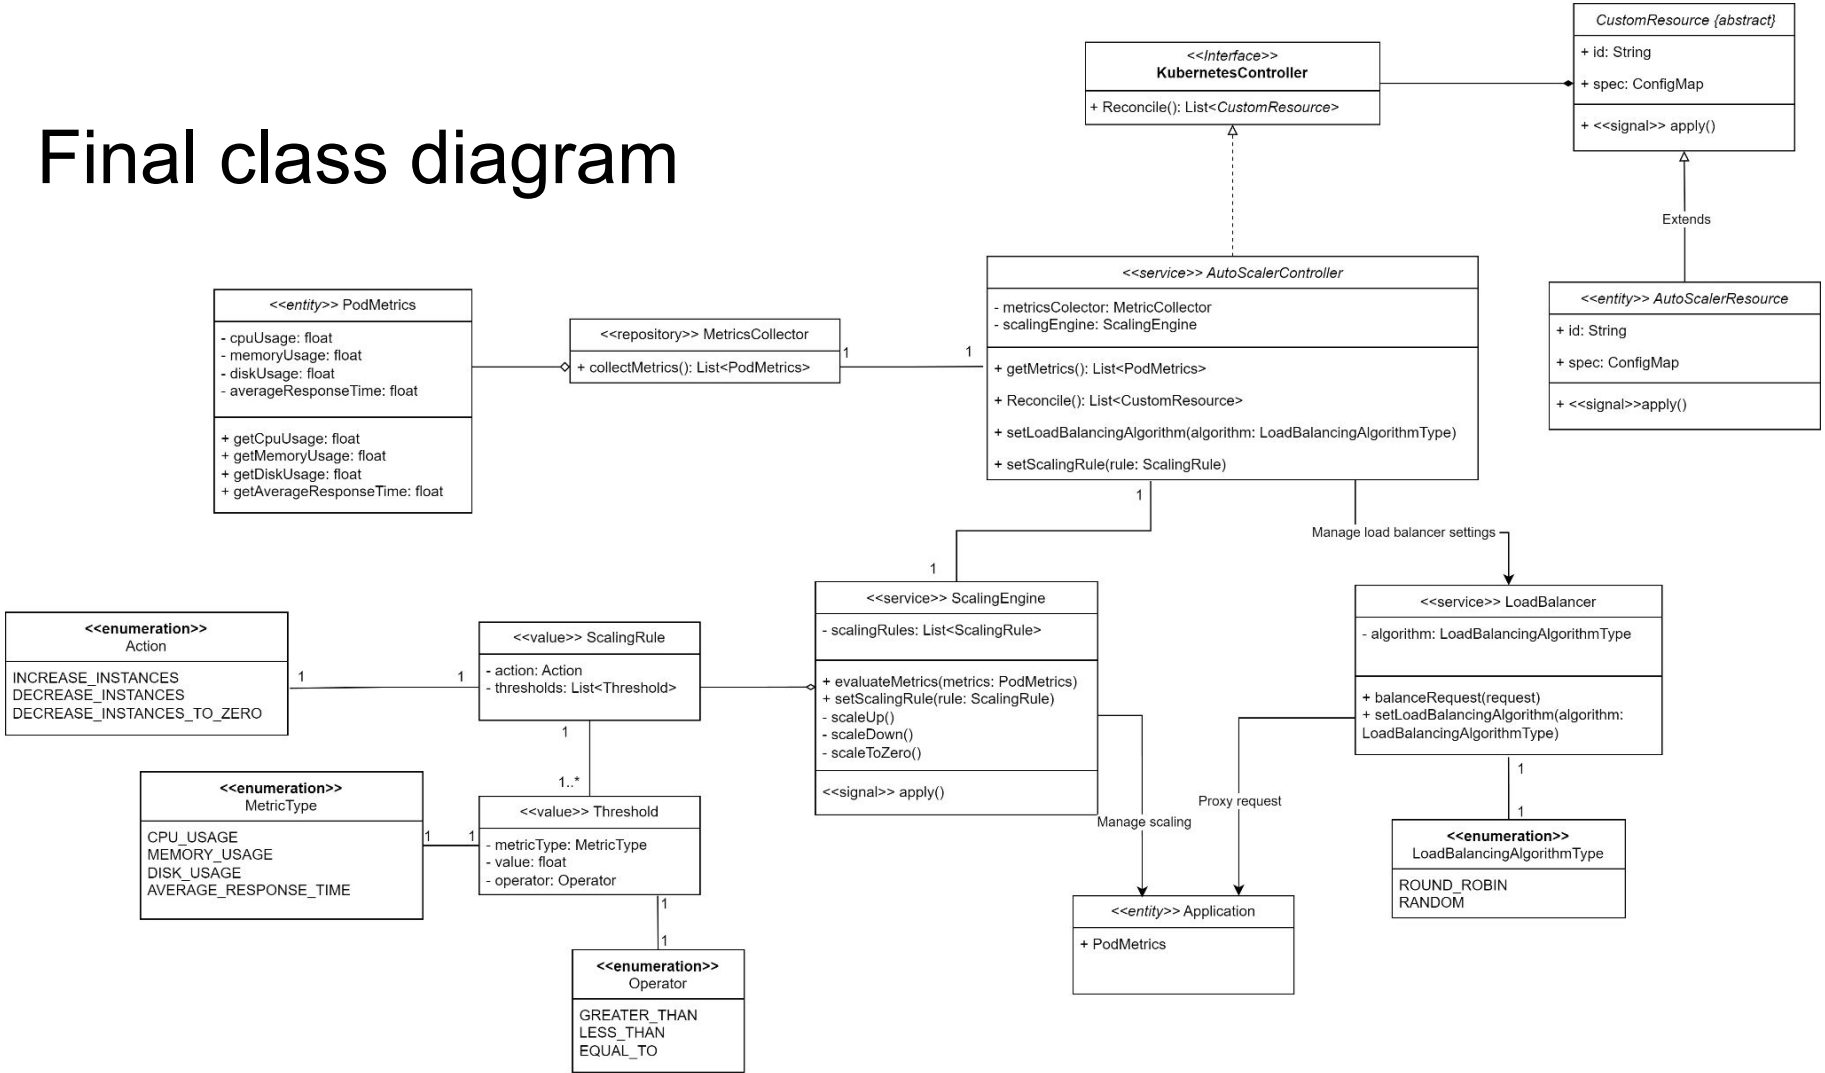
https://github.com/Lokhmat/k8s_nemesis/blob/main/final_task_materials/textual_use_cases.md



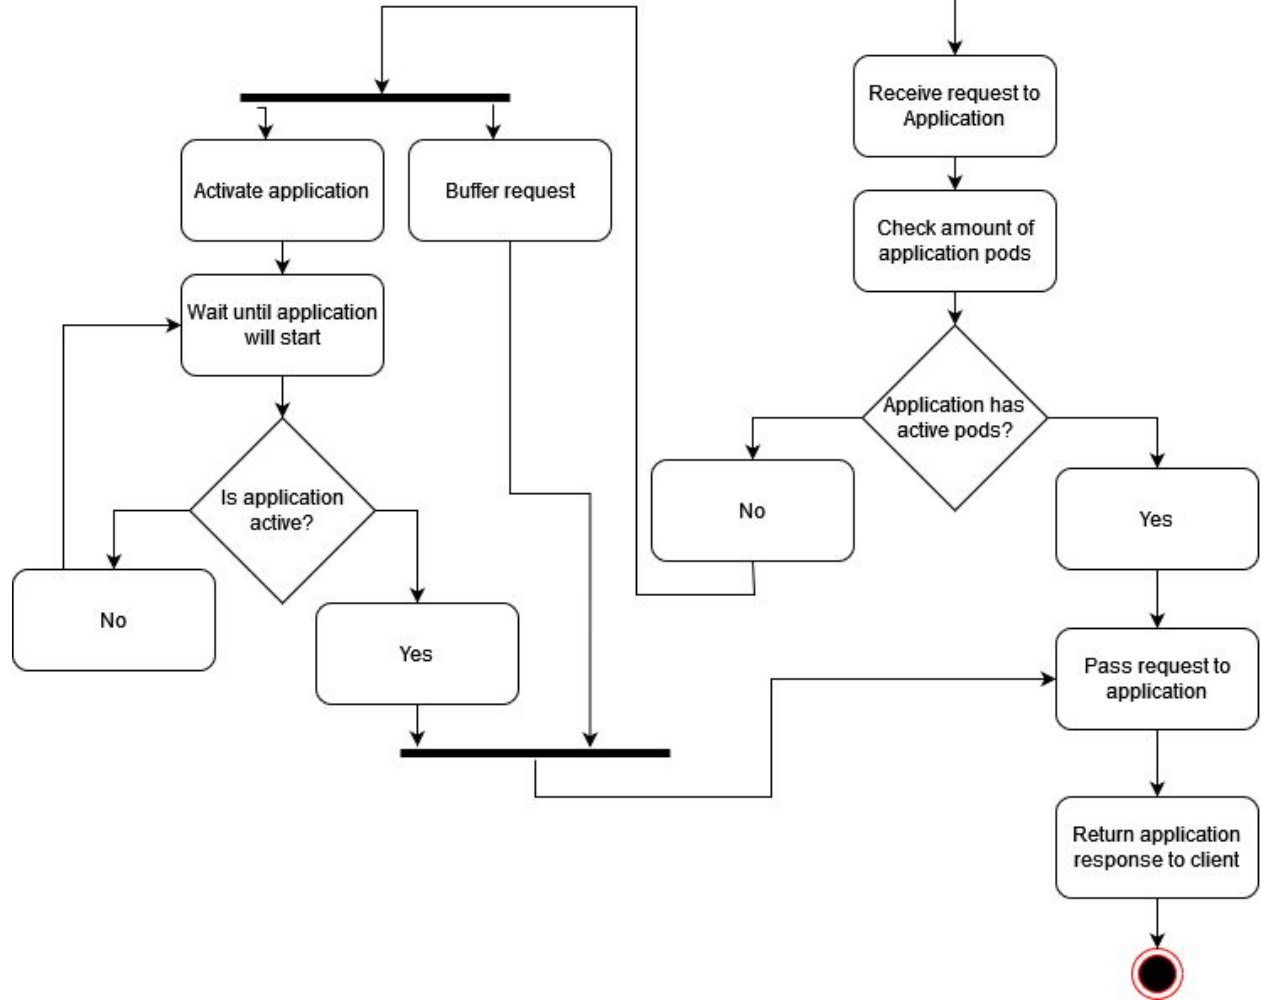
Interaction analysis

Use Case	Cooperation Type	Used Roles	Candidate Classes
Set scaling rule	Edit	User, ScalingEngine	ScalingEngine, ScalingRule, AutoScalerController, Treshold
Set load balancing policy	Edit	User, LoadBalancer	LoadBalancer, AutoScalerController
Set application disabling rule	Edit	User, ScalingEngine	ScalingEngine, ScalingRule, AutoScalerController, Treshold
Proxy request	Proxying	Application client, LoadBalancer, Application pod	LoadBalancer
Buffer request	Proxying	Application client, LoadBalancer	LoadBalancer
Activate cluster	Manage cluster	Application client, LoadBalancer, ScalingEngine	LoadBalancer
Get k8s cluster metrics	Validate	Kubernetes server, AutoScalerController	PodMetrics, MetricsCollector, AutoScalerController
Get current pods liveliness state	Validate	Kubernetes server, AutoScalerController, Application pod	PodMetrics, MetricsCollector, AutoScalerController
Decrease cluster pods amount	Manage cluster	Kubernetes server, AutoScalerController, Application pod	ScalingEngine, AutoScalerResource, AutoScalerController
Increase cluster pods amount	Manage cluster	Kubernetes server, AutoScalerController, Application pod	ScalingEngine, AutoScalerResource, AutoScalerController
Disable cluster	Manage cluster	Kubernetes server, AutoScalerController, Application pod	ScalingEngine, AutoScalerResource, AutoScalerController

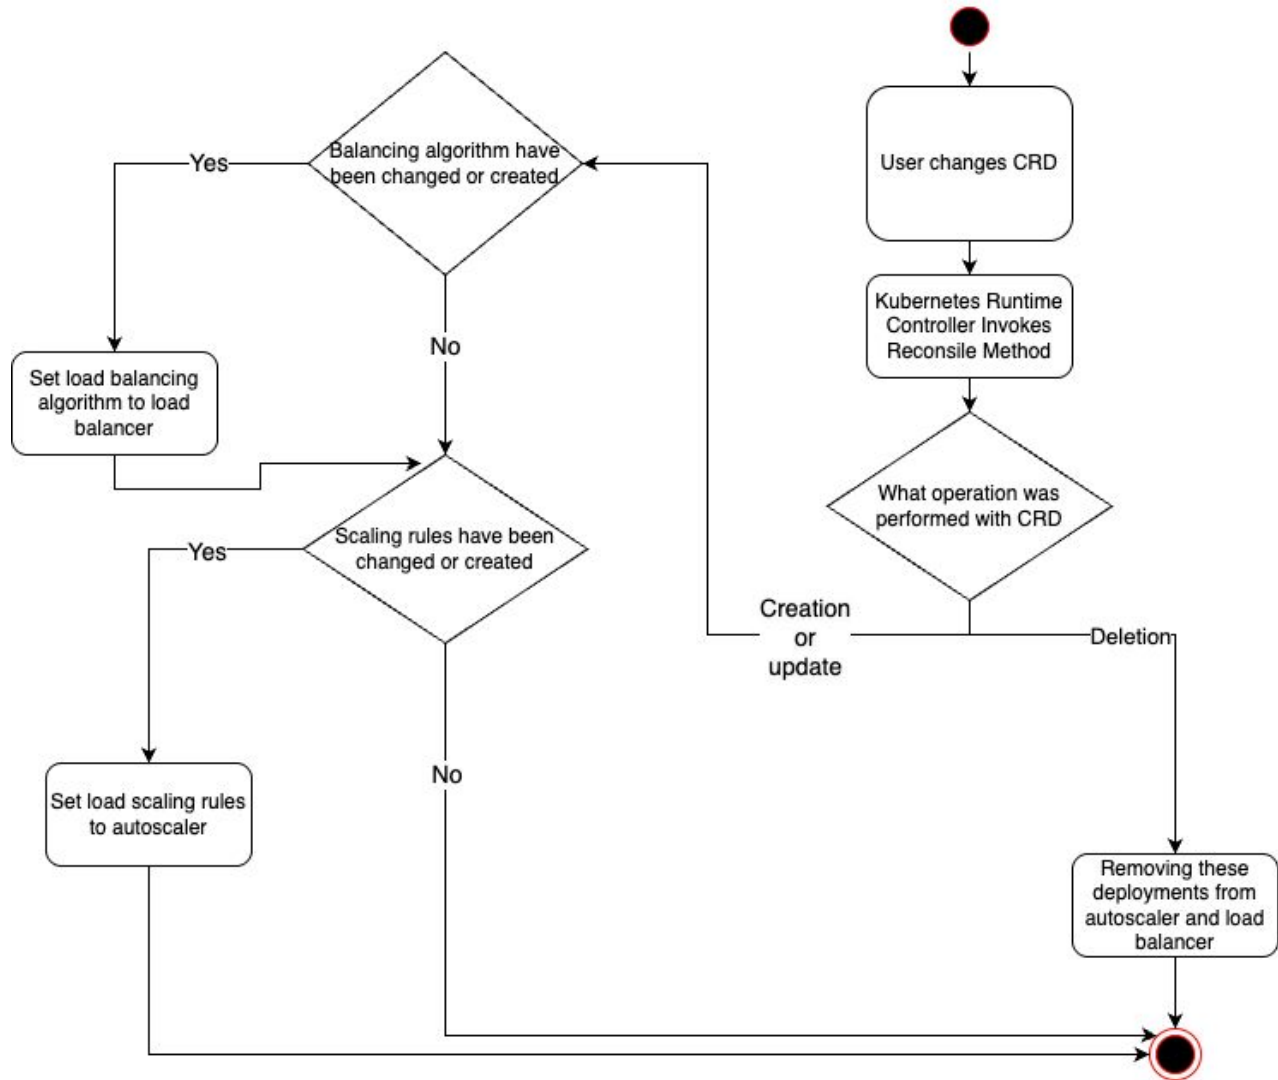
Final class diagram



Activity diagram

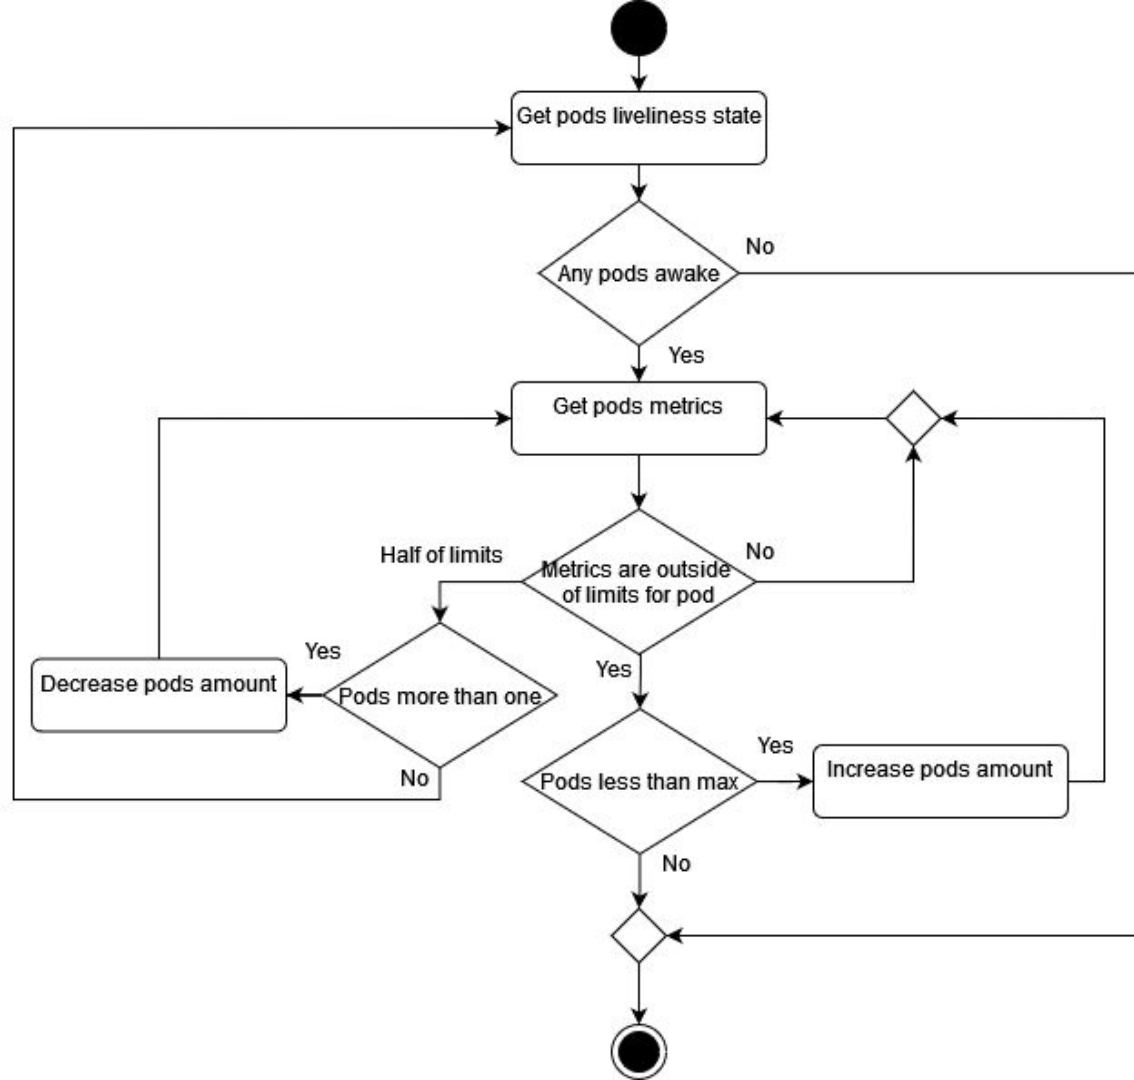


Configuration

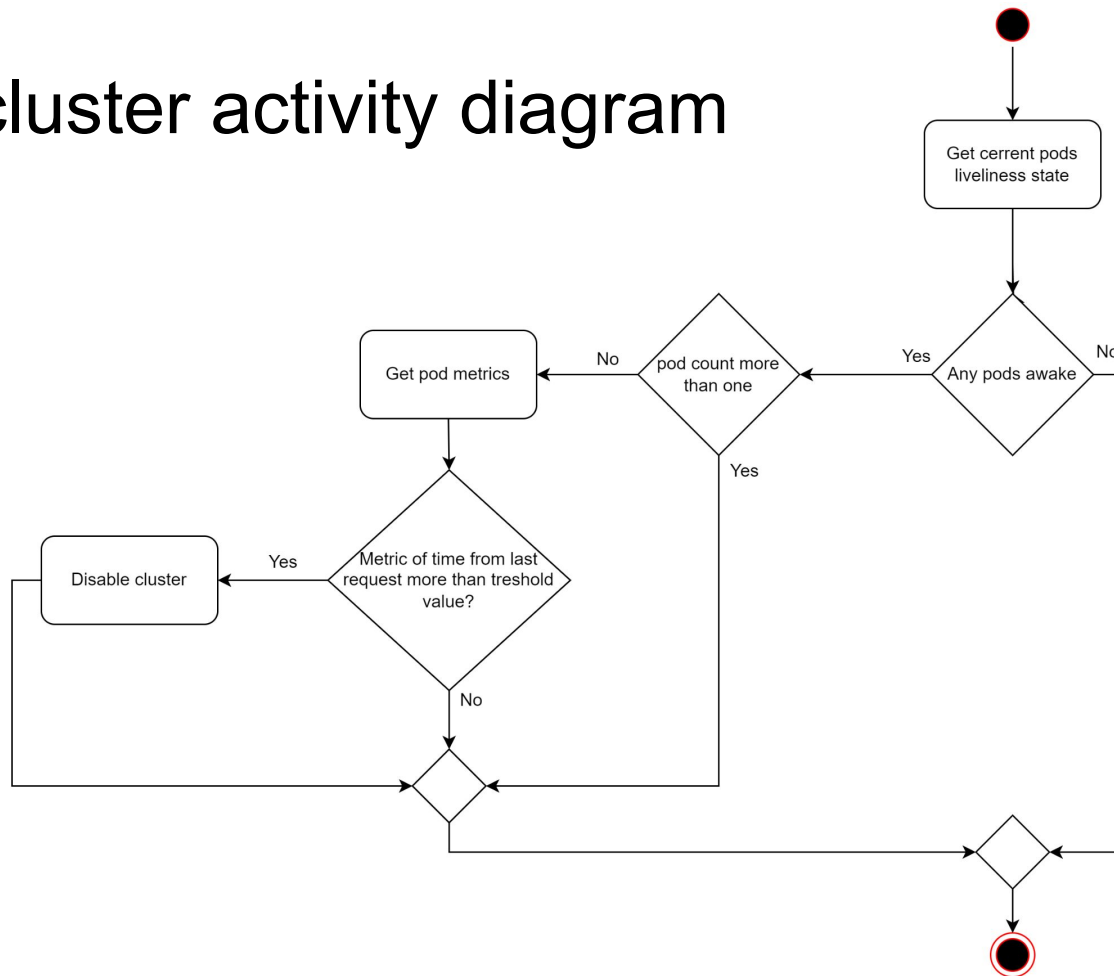


Decrease and Increase cluster pods amount

Activity diagram



Disable cluster activity diagram



Repository structure

main

1 Branch


Tags


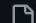







Go to file

t


Add file

<> Code

 **Lokhmat** Added textual use case description 1fb79d4 · 2 days ago 7 Commits

 final_task_materials	Added textual use case description	2 days ago
 CRC_cards.png	Added docs	2 weeks ago
 CRC_cooperation_table.png	Added docs	2 weeks ago
 CRC_know_do_interact.png	Added docs	2 weeks ago
 Task6_slides.pdf	Added slides for 6 task	4 days ago
 UML2_diagram_after_CRC.png	Added docs	2 weeks ago
 UML2_diagram_before_CRC.png	Added docs	2 weeks ago
 Use_case_diagram.png	Added docs	2 weeks ago
 project_description.md	Formatting	last month

Team and roles




Vasiliy Lopatkin
@DFCZok
Final class diagram
Interaction analysis



Sergey
Lokhmatikov
@Lohmat_Sergey
Story map
Use case diagram



Roman Kuzmenko
@definitely_not_rk
Personas
Behavioral diagram



Andrey Tamplon
@andreytamplon
Final class diagram
Interaction analysis